

Profiling

Profiling

The best way to make your application faster is by improving inefficient code. But how can you tell which code is inefficient?

You've got to measure it! Go has you covered with the built-in `pprof` toolset.

Types Of Profiling

- CPU Profiling
- Memory Profiling
- Blocking Profiling

CPU Profiling

- Most Common type of profiling
- Interrupts every 10ms and records the stack trace of currently running go routines

CPU profiling can be misleading as it doesn't take into account `latency`. This is due to the fact that the profiler is only taking samples of CPU usage of active go routines, and `sleeping` go routines won't be counted in the cpu profile.

Filippo Valsorda has a great video explanation here:

Fighting latency: the CPU profiler is not your ally

Memory Profiling

- Records allocations from stack trace
- By default records 1 in every 1,000 allocations (can be configured)
- Stack allocations are assumed to be free and not tracked
- Can be difficult to get a complete picture because it tracks allocations, and not use

Block Profiling

Block profiling records amount of time a go routine spent waiting for a shared resource (blocked).

Block profiling is typically used once you have feel you have eliminated cpu and memory performance issues.

A go routine is considered to be "blocked" when:

- Sending or receiving on a unbuffered channel.
- Sending to a full channel, receiving from an empty one.
- Trying to obtain a Lock on a sync.Mutex that is already locked. This can be a read or write lock.

One Profile At A Time

Profiling is not free. Profiling has a moderate, but measurable impact on program performance—especially if you increase the memory profile sample rate. Most tools will not stop you from enabling multiple profiles at once. If you enable multiple profiles at the same time, they will observe their own interactions and skew your results.

Do not enable more than one kind of profile at a time.

Rules Of Performance

1. Never guess about performance.
2. Measurements must be relevant.
3. Profile before you decide something is performance critical.
4. Test to know you are correct.

Using pprof

The pprof tool *should* always be invoked with *two* arguments.

```
$ go tool pprof /path/to/your/binary /path/to/your/profile
```

The "binary" argument *must* be the binary that produced this profile.

The "profile" argument *must* be the profile generated by this binary.

Warning: Because pprof also supports an online mode where it can fetch profiles from a running application over http, the pprof tool can be invoked without the name of your binary ([issue 10863](#)):

Do not do this or pprof will report your profile is empty:

```
$ go tool pprof /tmp/c.pprof
```

Generating A Profile From Benchmarks

To run a pprof against a benchmark, simply add the `-cpuprofile` flag as an argument.

In this example we will run a profile on a simple benchmark for the `strings.Split` function.

```
func BenchmarkSplit(b *testing.B) {  
    for i := 0; i < b.N; i++ {  
        split = strings.Split("one,two,three,four,five,six,seven", ",")  
    }  
}
```

Ø

Running The Benchmark

```
$ go test -bench=Split -cpuprofile=cpu.pprof
BenchmarkFields_Escape-8      5000000      244 ns/op
PASS
ok      github.com/gopherguides/training/advanced/profiling/src 1.475s
```

In addition to the benchmark output, two new files were created.

- `cpu.pprof`: The binary file that we will use with `pprof`
- `x.test`: where `x` corresponds to the name of the current directory. It's the test binary that runs the benchmarks.

Running The pprof Command

pprof is an interactive console. Typing help at any time will list off many of the things you can do, however, top is one of the more useful commands to run.

```
$ go tool pprof src.test cpu.pprof
Entering interactive mode (type "help" for commands)
(pprof) top
1.13s of 1.17s total (96.58%)
Showing top 10 nodes out of 45 (cum >= 0.61s)
```

flat	flat%	sum%	cum	cum%	
0.46s	39.32%	39.32%	0.46s	39.32%	runtime.mach_semaphore_signal
0.28s	23.93%	63.25%	1.10s	94.02%	strings.genSplit
0.13s	11.11%	74.36%	0.13s	11.11%	runtime.heapBitsSetType
0.12s	10.26%	84.62%	0.12s	10.26%	runtime.indexbytebody
0.06s	5.13%	89.74%	0.20s	17.09%	strings.Count
0.04s	3.42%	93.16%	0.04s	3.42%	runtime.mach_semaphore_wait
0.01s	0.85%	94.02%	0.01s	0.85%	nanotime
0.01s	0.85%	94.87%	0.01s	0.85%	runtime.mach_semaphore_timedwait
0.01s	0.85%	95.73%	0.62s	52.99%	runtime.makeslice
0.01s	0.85%	96.58%	0.61s	52.14%	runtime.mallocgc

Understanding Top

As you can see, 94% of our time was spent in `strings.genSplit`, if we look at the source code you will see it is just a wrapper for `strings.genSplit`:

```
// Split slices s into all substrings separated by sep and returns a slice of
// the substrings between those separators.
// If sep is empty, Split splits after each UTF-8 sequence.
// It is equivalent to SplitN with a count of -1.
func Split(s, sep string) []string { return genSplit(s, sep, 0, -1) }
```

Listing A Call

To investigate further, we can `list` the command:

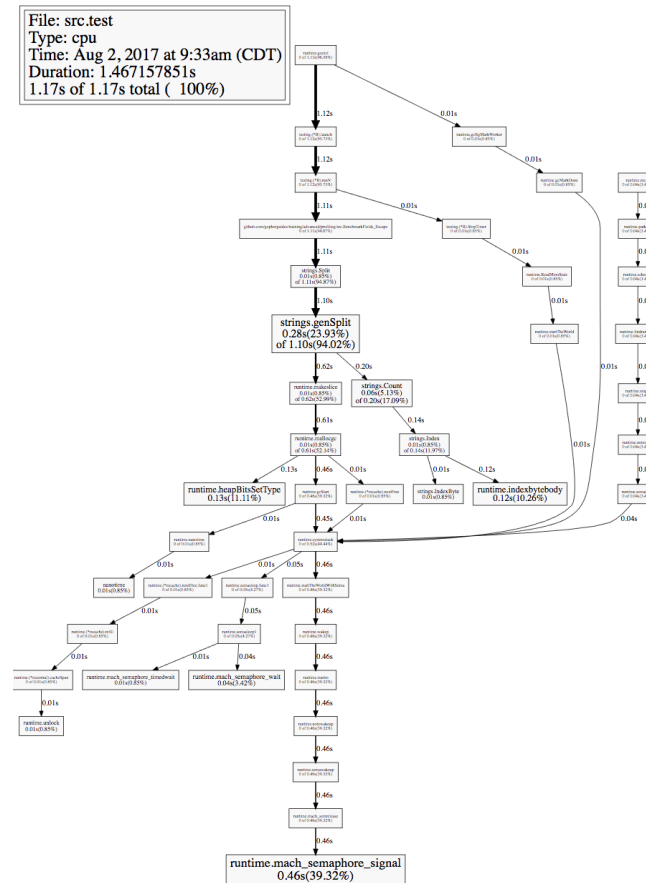
```
(pprof) list strings.genSplit
```

It's a large output, so here are a few lines from it:

```
ROUTINE ===== strings.genSplit in /usr/local/Cellar/go/1.8.3/libexec/src/strings
280ms      1.10s (flat, cum) 94.02% of Total
.          .      228:  return -1
.          .      229:}
.          .      230:
.          .      231:// Generic split: splits after each instance of sep,
.          .      232:// including sepSave bytes of sep in the subarrays.
10ms       10ms    233:func genSplit(s, sep string, sepSave, n int) []string {
10ms       10ms    234:  if n == 0 {
.          .      235:      return nil
.          .      236:  }
.          .      237:  if sep == "" {
.          .      238:      return explode(s, n)
.          .      239:  }
.          .      240:  if n < 0 {
.          200ms   241:      n = Count(s, sep) + 1
.          .      242:  }
.          .      243:  c := sep[0]
.          .      244:  start := 0
.          620ms   245:  a := make([]string, n)
.          .      246:  na := 0
```

Using Web

You can also see this in a visual format by using the `web` command. This will open up a browser and show visually where the time was spent.



Web Dependency

Using the web command required **Graphviz** to be installed on your machine. If you see this error you likely don't have it installed.

```
(pprof) web  
Failed to execute dot. Is Graphviz installed? Error: exec: "dot": executable file not found in $PATH
```


Running A Memory Profile (alloc)

The `--alloc_objects` flag tells you how many objects were allocated.

```
$ go test -bench=Split -memprofile=mem.pprof
BenchmarkSplit-8          5000000          239 ns/op
PASS
ok      github.com/gopherguides/training/advanced/profiling/src 1.455s
```

```
$ go tool pprof --alloc_objects src.test mem.pprof
Entering interactive mode (type "help" for commands)
(pprof) top
5997184 of 5997184 total ( 100%)
      flat flat% sum% cum cum%
5997184    100% 100% 5997184 100% strings.genSplit
      0      0% 100% 5997184 100% github.com/gopherguides/training/advanced/profiling/src
      0      0% 100% 5997184 100% runtime.goexit
      0      0% 100% 5997184 100% strings.Split
      0      0% 100% 5997184 100% testing.(*B).launch
      0      0% 100% 5997184 100% testing.(*B).runN
```

Running A Memory Profile (alloc) (cont.)

```
(pprof) list strings.genSplit
Total: 5997184
ROUTINE ===== strings.genSplit in /usr/local/Cellar/go/1.8.3/libexec/src/strings
5997184      5997184 (flat, cum) 100% of Total
.           .      240:  if n < 0 {
.           .      241:      n = Count(s, sep) + 1
.           .      242:  }
.           .      243:  c := sep[0]
.           .      244:  start := 0
5997184      5997184  245:  a := make([]string, n)
.           .      246:  na := 0
.           .      247:  for i := 0; i+len(sep) <= len(s) && na+1 < n; i++ {
.           .      248:      if s[i] == c && (len(sep) == 1 || s[i:i+len(sep)] == sep) {
.           .      249:          a[na] = s[start : i+sepSave]
.           .      250:          na++
```

Running A Memory Profile (alloc)

The `--inuse_objects` flag tells you how many objects are still in use.

```
$ go tool pprof --inuse_objects src.test mem.pprof
Entering interactive mode (type "help" for commands)
(pprof) top
0 of 0 total ( 0%)
      flat flat%   sum%        cum      cum%
      0      0%     0%          0         0%  github.com/gopherguides/training/advanced/profiling/src
      0      0%     0%          0         0%  runtime.goexit
      0      0%     0%          0         0%  strings.Split
      0      0%     0%          0         0%  strings.genSplit
      0      0%     0%          0         0%  testing.(*B).launch
      0      0%     0%          0         0%  testing.(*B).runN
```

Due to the benchmark being very small, we actually have no measurable in use objects. We will revisit this command later in the module.

Exercise

- Write a benchmark for the `strings.ToUpper` function.
- Uppercase the following string in the benchmark: Education is what remains after one has forgotten what one has learned in school.
- Using the pprof tool, generate a cpu profile
- Use the following commands: `top`, `list`, `web`
- Extra Credit - Name who the quote is from.

Solution

```
$ go test -bench=ToUpper -cpuprofile=cpu.pprof
BenchmarkToUpper-8      20000000      858 ns/op
PASS
ok      github.com/gopherguides/training/advanced/profiling/src 2.582s
```

```
package profiling

import (
    "strings"
    "testing"
)

var upper string

func BenchmarkToUpper(b *testing.B) {
    for i := 0; i < b.N; i++ {
        upper = strings.ToUpper("Education is what remains after one has forgotten what one has learned")
    }
}
```

Profiling Live Programs

Using `net/http/pprof`, we can grab live profiles from running programs via an http endpoint.

You will need to import the package with an underscore (`_`), since we are importing it for side effects, and not direct use.

```
import _ "net/http/pprof"
```

Once you have the program running, you can visit the <http://localhost:8080/debug/pprof/> and explore the dashboard.

Demo App

```
// super simple email regex
const emailRegex = `(\w+)\@(\w+)\. [a-zA-Z]`

func main() {
    http.HandleFunc("/", handler)
    log.Printf("listening on localhost:8080")
    log.Fatal(http.ListenAndServe("localhost:8080", nil))
}

func handler(w http.ResponseWriter, r *http.Request) {
    email := r.URL.Path[1:]

    if isEmail(email) {
        log.Printf("%s looks like an email\n", email)
        fmt.Fprintf(w, "%s looks like an email", email)
        return
    }
    log.Printf("%s doesn't look like an email\n", email)
    fmt.Fprintf(w, "%s doesn't look an email", email)
}

func isEmail(email string) bool {
    re := regexp.MustCompile(emailRegex)
    return re.MatchString(email)
}
```



Grabbing A Live Profile

```
$ go tool pprof -seconds 5 http://localhost:8080/debug/pprof/profile  
Fetching profile from http://localhost:8080/debug/pprof/profile?seconds=5  
Please wait... (5s)  
Saved profile in /Users/corylanou/pprof/pprof.localhost:8080.samples.cpu.001.pb.gz  
Entering interactive mode (type "help" for commands)  
(pprof)
```


The Profile

```
(pprof) top  
profile is empty
```

What happened?

When we were running the profile, there was no actual load on the server, so nothing was recorded

Adding Load

We can add some load to the server with the **go-wrk** tool.

```
go get github.com/adjust/go-wrk
```

Creating load:

```
go-wrk -d 30 http://localhost:8080/john@smith.com
```

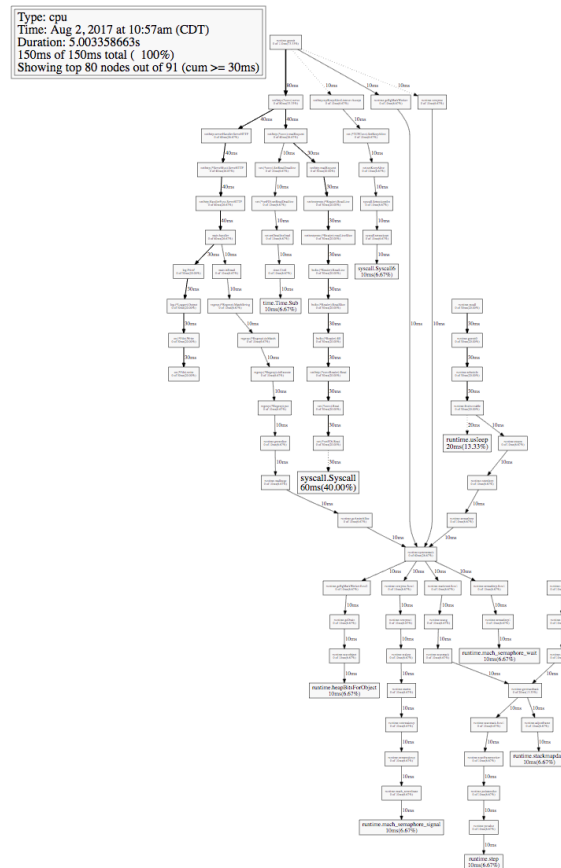
Now, while that is working, we can grab another profile:

```
$ go tool pprof -seconds 5 http://localhost:8080/debug/pprof/profile
```


Let's use the **web** command to look at the results.

Inspecting Results Via Web

Wow, this is a lot of information, and hard to follow. Let's move on to flame graphs and see if that helps us understand better what is going on.



Flame Graphs

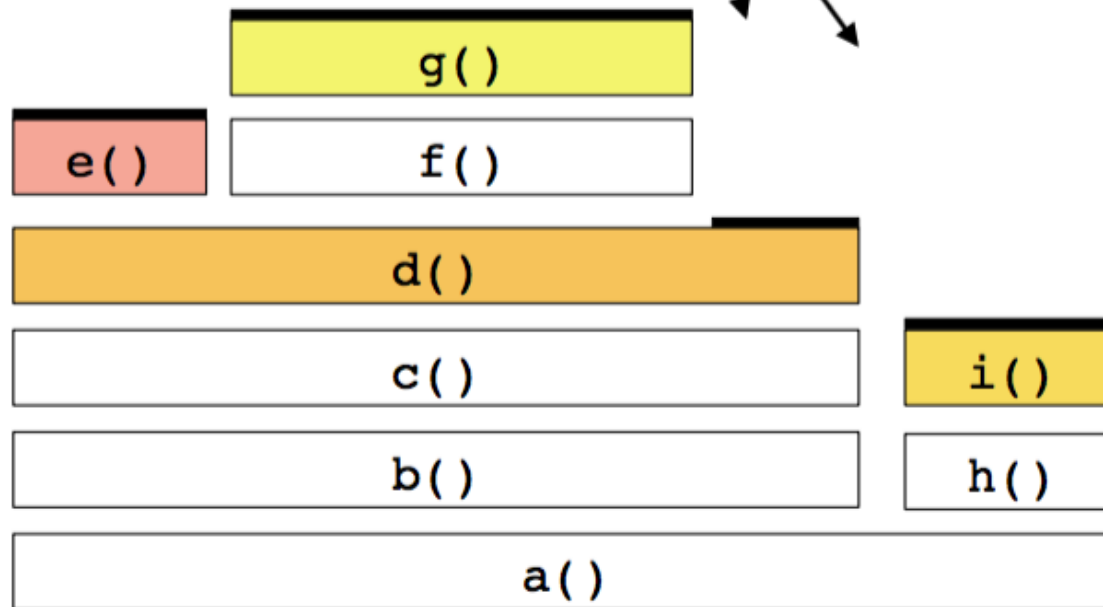


"The x-axis shows the stack profile population, sorted alphabetically (it is not the passage of time), and the y-axis shows stack depth. Each rectangle represents a stack frame. The wider a frame is, the more often it was present in the stacks. The top edge shows what is on-CPU, and beneath it is its ancestry. The colors are usually not significant, picked randomly to differentiate frames."

- *x axis* alphabetical stack sort, to maximize merging.
- *y axis* stack depth.

Flame Graph Interpretation (1/3)

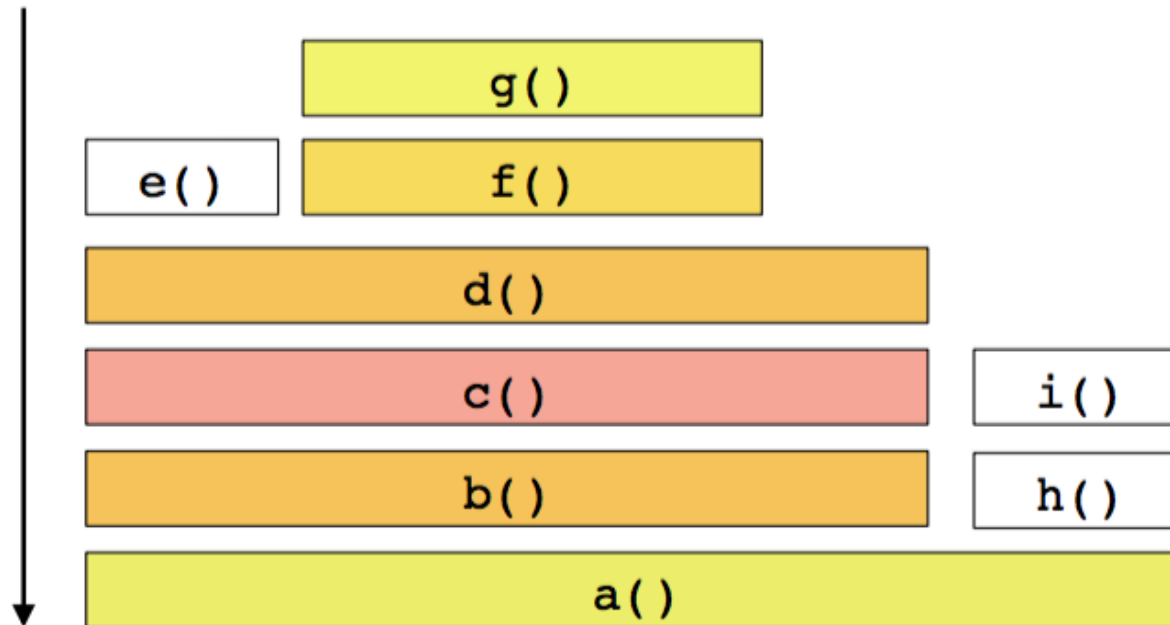
Top edge shows who is running on-CPU,
and how much (width)



Flame Graph Interpretation (2/3)

Top-down shows ancestry

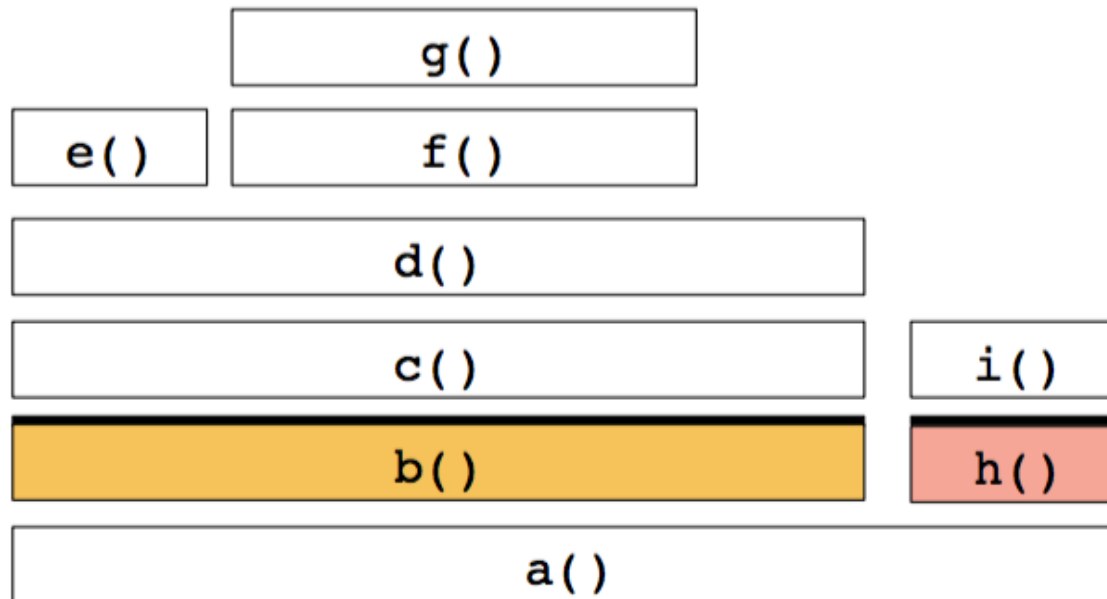
e.g., from g():



Flame Graph Interpretation (3/3)

Widths are proportional to presence in samples

e.g., comparing b() to h() (incl. children)



Using Go-torch

Flame graphs can consume data from many sources, including `pprof` (and `perf(1)`).

Uber has open sourced a tool call `go-torch` which automates the process, assuming you have the `/debug/pprof` endpoint, or you can feed it an existing profile.

DEMO

Get needed support scripts

```
$ go get -u -v github.com/uber/go-torch
$ mkdir -p $GOPATH/src/github.com/brendangregg
$ cd $GOPATH/src/github.com/brendangregg
$ git clone https://github.com/brendangregg/FlameGraph.git
$ export PATH=$PATH:$GOPATH/src/github.com/brendangregg/FlameGraph
```

Run the torch command:

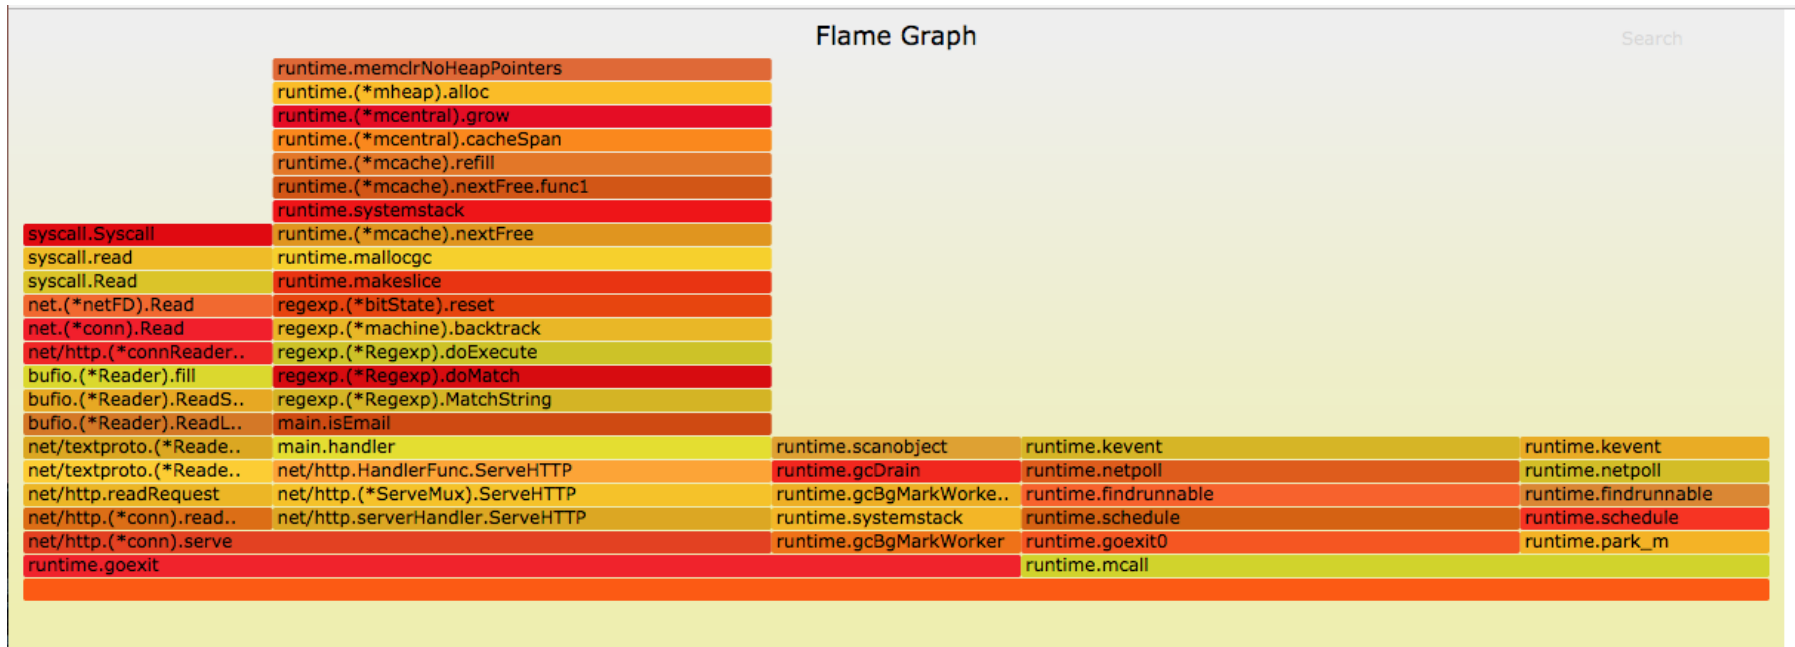
```
$ go-torch -t 5

INFO[11:14:35] Run pprof command: go tool pprof -raw -seconds 5 http://localhost:8080/debug/pprof/p
INFO[11:14:41] Writing svg to torch.svg
```

Inspecting Torch Output

The width represents relative time in a call. The wider the column, the more time spent in that call.


The colors and the order on the graph are random and hold no value other than making it easy to distinguish one call from another.



Searching Flame Graphs

There is a `search` input box that is very hard to see in the upper right hand side.

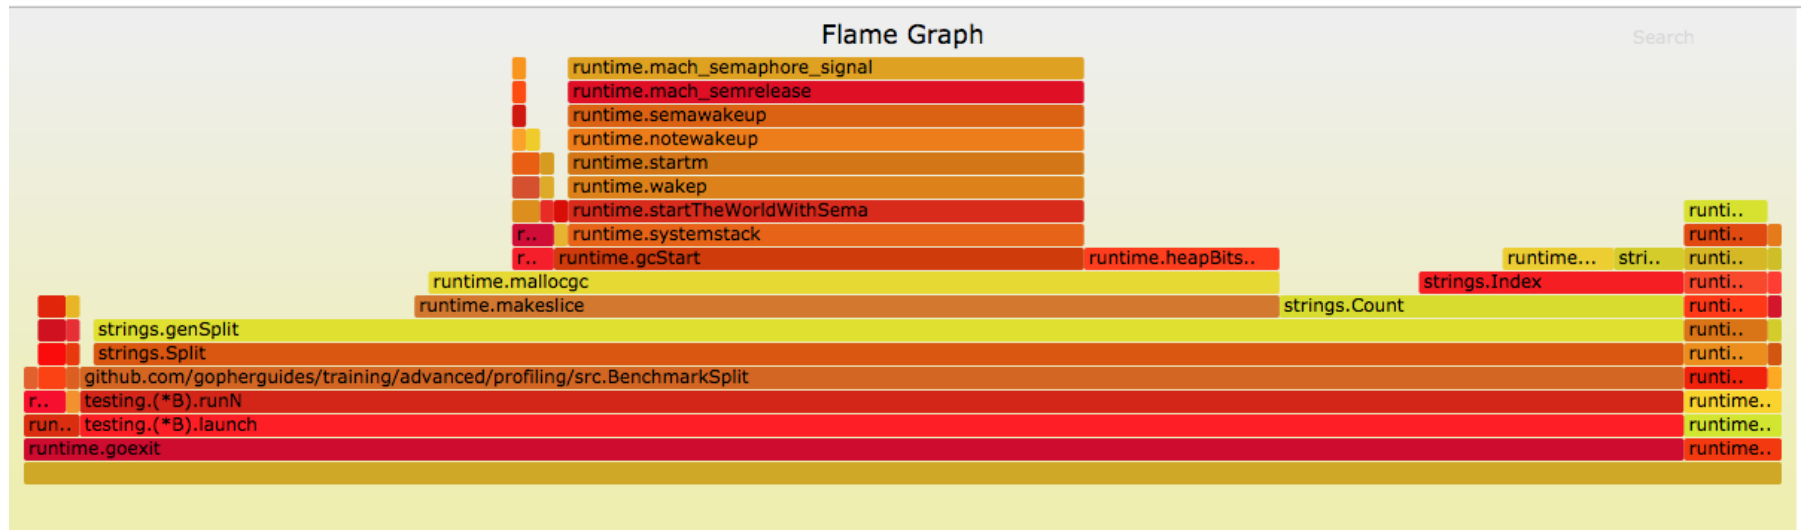
Using it we can search for the `isEmail` function and highlight that part of the flame graph.



Generating From Existing Profiles

We can also use the output from our benchmark profile to generate flame graphs as well.

```
$ go-torch --binaryname src.test -b cpu.pprof
INFO[11:24:37] Run pprof command: go tool pprof -raw src.test cpu.pprof
INFO[11:24:37] Writing svg to torch.svg
```



Exercise

- Using the output of the benchmarks on the previous `ToUpper` exercise, generate a flame graph.